

Gated Position-Aided Fusion for Blockage-Robust Neural Beam Selection in 5G NR mmWave Systems

Furkan Yakkan and Akif Hasdemir
Graduate School of Engineering and Science
Abdullah Gül University, Kayseri, Türkiye
Email: {furkan.yakkan, akif.hasdemir}@agu.edu.tr

Abstract—Beam selection in millimeter-wave (mmWave) 5G New Radio (NR) requires sweeping a large beam-pair codebook, incurring substantial initial-access overhead. A well-known neural approach predicts the full reference-signal-received-power (RSRP) profile over all beam pairs from a small set of downsampled RSRP measurements, enabling a top- K sweep over only the most promising beams. We first reproduce this regression benchmark on the standardized TR 38.901 urban macrocell scenario, attaining 90% top- K accuracy at $K = 13$ (81.4% sweep-overhead reduction). We then investigate whether fusing the already available user-equipment (UE) position with the sparse RSRP improves selection. Contrary to the common multi-modal intuition, we find that on *clean* RSRP the single-modality predictor is already near-optimal: position fusion, learned beam sampling, a beam-grid convolutional reconstruction, and uncertainty-aware adaptive- K all fail to beat it. The value of the position modality emerges only when the RSRP modality is *degraded*. Building on this, we propose a *gated residual position-aided fusion* network in which a position-driven correction is added to the RSRP-based estimate through a learned gate. The gate preserves clean-data accuracy at no cost while the residual path delivers a consistent 2–4 percentage-point top- K accuracy gain under mmWave beam blockage (when 6–12 of the 14 measured beams are lost). An architecture ablation confirms the gate is responsible for the no-regret behavior. The result is a practical, no-regret robustness improvement for blockage-prone mmWave links.

Index Terms—5G NR, millimeter wave, beam management, beam selection, deep learning, multi-modal fusion, beam blockage, RSRP.

I. INTRODUCTION

Millimeter-wave (mmWave) bands are central to 5G NR and 6G for their large bandwidths, but their high pathloss and susceptibility to blockage demand highly directional beamforming [1]. During initial access, the gNB and UE must establish a good transmit/receive beam pair, classically through an exhaustive synchronization-signal-block (SSB) beam sweep over the full codebook. With many narrow beams, this sweep dominates initial-access latency and overhead, motivating machine-learning methods that predict good beams from partial or out-of-band observations [2], [3].

Two broad families have emerged. *In-band* methods exploit RSRP or channel measurements [4], [5]; they are accurate when measurements are reliable. *Side-information* methods exploit position, LiDAR, radar, vision, or sub-6 GHz channels [6]–[10]; they are robust to RF impairments but typically

achieve lower peak accuracy. A widely used in-band formulation, distributed as a reference example by MathWorks, casts beam selection as RSRP-profile *regression*: a neural network maps a downsampled subset of RSRP measurements to the full per-beam-pair RSRP profile, after which only the top- K predicted beams are swept.

In this work we take this regression benchmark as our starting point and ask a simple question: *does adding the UE position—already available at the network—improve beam selection?* Our investigation yields a nuanced and, we believe, instructive answer. The contributions are:

- We reproduce the SSB-RSRP regression benchmark on the TR 38.901 urban macrocell (UMa) scenario and confirm its headline operating point: 90% top- K accuracy at $K = 13$, i.e. an 81.4% reduction in beam-sweep overhead, with average RSRP within 0.24 dB of exhaustive search (Section III).
- We show that, on clean RSRP, the single-modality predictor is at the achievable ceiling: naive position fusion, statistically “optimized” beam sampling, a 2D beam-grid convolutional reconstruction, and uncertainty-aware adaptive- K all fail to improve it (Section IV). This isolates *where* a second modality can help.
- We propose a *gated residual position-aided fusion* network and an impairment-aware training procedure (Section V). It is a *no-regret* design: it matches the RSRP-only model on clean data and improves top- K accuracy by 2–4 points under beam blockage. An architecture ablation shows the gate is what preserves the clean-data parity (Section VI).

II. RELATED WORK

Early ML beam-selection work posed the task as classification of the best beam (pair) from side information such as UE position or LiDAR [2], [3], [6]. Position/location-aided prediction was extended to 3D beam selection with orientation information [7], and to real-world GPS-based prediction on the DeepSense 6G testbed [8]. Out-of-band approaches predict mmWave beams (and blockages) from sub-6 GHz channels [4], [10], [11], reporting large overhead reductions (e.g. $\sim 79\%$ in [4]). Sensing-aided prediction using radar [9] and LiDAR [12] has been demonstrated on real V2I data, and waveform-learning approaches such as DeepBeam [5] reduce

beam-management latency without coordination. A parallel line predicts *blockage* for proactive handoff [13], [14], and recent work exploits temporal correlation for vehicular beam management [15]. Surveys [1] and the 3GPP study on AI/ML for the NR air interface [16] provide broader context.

Table I summarizes representative methods by input modality and reported benefit. Two gaps motivate our work. First, multi-modal fusion is typically studied with rich external sensors (camera/LiDAR/radar) on vehicular data, not with the *already-available* UE position fused into the canonical SSB-RSRP regression. Second, prior fusion work emphasizes peak accuracy on nominal data rather than *robustness as a function of measurement reliability*. We address both, and—importantly—report the negative result that on clean RSRP the position modality is redundant, which clarifies that the fusion payoff is a robustness, not an accuracy, story.

III. SYSTEM MODEL AND REGRESSION BENCHMARK

A. Scenario and beam codebook

We adopt the TR 38.901 UMa scenario [17] in frequency range 2 (FR2) at a carrier of 30 GHz, following the AI/ML system-level assumptions of TR 38.843 [16]. The gNB uses a 4×8 cross-polarized panel and the UE a 1×4 panel, yielding $N_{\text{tx}} = 10$ transmit beams and $N_{\text{rx}} = 7$ receive beams, for a total of $N_b = N_{\text{tx}}N_{\text{rx}} = 70$ beam pairs covering a 120° sector. For each UE location, SSB-based beam sweeping yields an RSRP value for every beam pair, forming an RSRP matrix $\mathbf{R} \in \mathbb{R}^{7 \times 10}$; we index beam pairs by the column-major linear index $b \in \{1, \dots, 70\}$ of $\mathbf{R}$. We use the official pre-recorded dataset of 15,000 training and 500 test UE locations. Blocked beams appear as $-\infty$ RSRP (4.6% of entries; $\sim 10\%$ of locations have at least one blocked beam among the sampled set); we floor them to -120 dB for network inputs/targets.

B. Regression formulation

Let $\mathbf{r} \in \mathbb{R}^{70}$ be the vectorized RSRP profile of a UE. A fixed downsampling selects every fifth beam pair, giving $\mathcal{S} = \{1, 6, 11, \dots, 66\}$ with $|\mathcal{S}| = 14$ measured beams and input $\mathbf{x} = \mathbf{r}_{\mathcal{S}} \in \mathbb{R}^{14}$. A network $f_\theta : \mathbb{R}^{14} \rightarrow \mathbb{R}^{70}$ regresses the full profile $\hat{\mathbf{r}} = f_\theta(\mathbf{x})$; targets are min-max normalized to $[-1, 1]$ and the output layer is $\tanh$. At inference, the K beams with the largest predicted RSRP are swept and the one with the strongest *actual* RSRP is selected. The metrics are the top- K accuracy (the fraction of UEs whose true-optimal beam is among the K predicted) and the average selected RSRP.

The benchmark network is a four-hidden-layer multilayer perceptron (MLP) of width 96 with leaky-ReLU activations and z -score input normalization, trained with Adam to minimize mean-squared error. Position is *not* an input to the network; it is used only by a K -nearest-neighbor (KNN) benchmark, alongside Statistical and Random baselines.

C. Reproduction results

Table II reports the reproduced metrics and Fig. 1 the top- K accuracy curves. The neural predictor reaches 90.0% top- K accuracy at $K = 13$ —sweeping fewer than one quarter of




Fig. 1. Reproduced benchmark: top- K accuracy of the RSRP-regression network vs. KNN, Statistical, and Random baselines. The network reaches 90% at $K = 13$.

the 70 beam pairs, an 81.4% overhead reduction—and 97.4% at $K = 30$. Its average RSRP at $K = 13$ is -23.21 dB, within 0.24 dB of the exhaustive optimum (-22.97 dB) and well above KNN (-25.29 dB). The neural method dominates KNN, Statistical, and Random across all K , reproducing the published behavior and validating our pipeline.

IV. WHERE A SECOND MODALITY DOES (NOT) HELP

Before designing a fusion model, we tested whether the benchmark can be improved on *clean* RSRP along four axes, each under identical data, splits, and seeds:

- 1) **Naive position fusion**: a two-branch network ingesting the 14 RSRP values and the 3D UE position. Result: a tie within noise on clean data ($+0.8$ at $K=13$, -3.0 at $K=5$).
- 2) **Learned/statistical beam sampling**: replacing the uniform every-fifth sampling with the 14 highest-variance, most-frequently-optimal, or strongest-mean beams. All were *worse* (e.g. $K_{90}=39$ vs. 13), because clustered beams lose spatial coverage; uniform sampling is near-optimal.
- 3) **2D beam-grid CNN**: treating sparse sensing as inpainting on the 7×10 ($R_x \times T_x$) RSRP image with observed-value and mask channels. Result: worse at the operating point ($K_{90}=16$ vs. 13).
- 4) **Uncertainty-aware adaptive- K** : a nucleus rule over the predicted profile to choose a per-UE K . Result: negligible ($\leq +0.6$ pts at matched mean overhead).

The common explanation is that the optimal beam is determined by the channel, which the RSRP directly measures; the UE position predicts the optimal beam only weakly (KNN attains 37.8% at $K=13$). Hence, when RSRP is clean, position is largely *redundant*, and richer architectures merely add variance. This motivates targeting the regime where RSRP is unreliable.

TABLE I
REPRESENTATIVE ML BEAM-MANAGEMENT METHODS (2018–2025).

Work (year)	Modality	Approach	Reported benefit
Heng & Andrews [2] (2019)	Position	NN beam alignment	ML-assisted alignment vs. exhaustive sweep
Klautau <i>et al.</i> [6] (2019)	LiDAR	DL beam selection	Out-of-band beam selection
Sim <i>et al.</i> [4] (2020)	Sub-6 GHz	DL beam selection (NR/6G)	~79% sweep-overhead reduction
Alrabeiah & Alkhateeb [11] (2020)	Sub-6 GHz	DL beam + blockage prediction	>90% blockage prediction
Rezaie <i>et al.</i> [7] (2022)	Position + orientation	DL 3D beam set selection	Location/orientation-aided selection
Polese <i>et al.</i> [5] (2021)	Waveform/RF	Coordination-free beam mgmt	Up to 7× lower latency
Demirhan & Alkhateeb [9] (2022)	Radar	DL beam prediction	Real-world radar-aided prediction
Jiang <i>et al.</i> [12] (2023)	LiDAR	DL future-beam prediction	~90% overhead reduction (V2I)
Wu <i>et al.</i> [14] (2022)	In-band power	Moving-blockage prediction	>85% blockage accuracy
Oliveira <i>et al.</i> [15] (2025)	Position + time	RNN temporal beam mgmt	Up to 50% overhead reduction
This work	Sparse RSRP + position	Gated residual fusion	No clean-data cost; +2–4 pts under blockage

TABLE II
REPRODUCED BENCHMARK METRICS (15K/500, $K = 13$ UNLESS NOTED).

Metric	Neural	KNN	Optimal
Top- K reaches 90% at	$K = 13$	–	–
Overhead reduction @ $K=13$	81.4%	–	–
Top- K accuracy @ $K=13$	90.0%	37.8%	–
Top- K accuracy @ $K=30$	97.4%	–	–
Avg. RSRP @ $K=13$ (dB)	−23.21	−25.29	−22.97

V. PROPOSED METHOD: GATED RESIDUAL POSITION-AIDED FUSION

A. Beam-blockage model

Practical mmWave sweeps lose measurements to blockage and missed detections. We model this by, for each UE, setting n_B of the 14 measured beams (chosen uniformly at random) to the floor value, plus a small Gaussian RSRP measurement noise. The number of blocked input beams n_B parameterizes the impairment severity; n_B large corresponds to only a handful of reliably measured beams.

B. Architecture

The proposed network has two encoders. An RSRP encoder maps $\mathbf{x}$ to a primary estimate $\mathbf{b} = g_{\text{base}}(\mathbf{x}) \in \mathbb{R}^{70}$ (structurally the RSRP-only path). A position encoder maps the 3D UE position $\mathbf{p}$ to a correction $\mathbf{c} = g_{\text{corr}}(\mathbf{p}) \in \mathbb{R}^{70}$. A gate $\mathbf{g} = \sigma(g_{\text{gate}}([\mathbf{h}_r, \mathbf{h}_p])) \in (0, 1)^{70}$, computed from both encoder features $\mathbf{h}_r, \mathbf{h}_p$, weights the correction. The output is

$$\hat{\mathbf{r}} = \tanh(\mathbf{b} + \mathbf{g} \odot \mathbf{c}), \quad (1)$$

where $\odot$ is the elementwise product. The residual-plus-gate structure lets the network recover the strong RSRP-only estimate (by suppressing $\mathbf{g} \odot \mathbf{c}$) when RSRP is reliable, and inject the position prior when it is not. The RSRP encoder uses two width-96 layers, the position encoder two width-32 layers, and the gate a width-64 layer; all use leaky-ReLU.

TABLE III
TOP- K ACCURACY ($K=13$, %) VS. NUMBER OF BLOCKED INPUT BEAMS.

# blocked	0	2	4	6	8	10	12
RSRP-only	89.2	84.6	78.1	69.8	56.0	35.9	22.3
Gated fusion	89.4	83.5	78.7	72.1	57.9	39.7	24.9
Gain (pts)	+0.2	−1.1	+0.6	+2.3	+2.0	+3.8	+2.6

C. Impairment-aware training

Both the proposed model and the RSRP-only reference are trained with the *same* blockage augmentation: per sample, $n_B \sim \mathcal{U}\{0, \dots, 8\}$ beams are dropped and light noise is added. This teaches the models to operate across blockage levels and makes the comparison differ only in the position modality. Networks are trained with Adam (300 epochs, mini-batch 512, initial learning rate 10^{-3} with piecewise decay).

VI. EXPERIMENTAL RESULTS

A. Robustness to beam blockage

Fig. 2 and Table III report top- K accuracy ($K = 13$, averaged over 10 random blockage draws) versus the number of blocked input beams. On clean inputs ($n_B = 0$) the proposed fusion matches the RSRP-only model (89.4% vs. 89.2%): the position modality incurs *no* cost. As beams are blocked, the fusion pulls ahead, with a consistent +2 to +3.8-point advantage for $n_B \in [6, 12]$ (e.g. 39.7% vs. 35.9% at $n_B = 10$). The position-only KNN is blockage-immune but flat at 37.8%, far below both networks except under extreme blockage. Fig. 3 shows that, at a representative $n_B = 8$, the fusion dominates the RSRP-only model across the entire K range.

B. Architecture ablation: is the gate needed?

Table IV compares four fusion variants trained identically: RSRP-only, concatenation fusion, plain residual ($\mathbf{g} = 1$ in (1)), and the proposed gated residual. The plain residual already provides most of the blockage robustness but costs 1.2 points




Fig. 2. Top- K accuracy ($K=13$) vs. number of blocked input beams. The gated fusion matches RSRP-only when clean and improves it under blockage; KNN (position only) is flat.




Fig. 3. Top- K accuracy vs. K at a representative 8/14 blocked input beams. The gated fusion dominates the RSRP-only model across K ; both far exceed position-only KNN.

on clean data (88.0% vs. 89.2%). The gate restores clean parity (89.4%) at equal blockage robustness—this is precisely the role of the gate. Interestingly, ungated *concatenation* achieves the largest heavy-blockage gains (45.5% vs. 36.2% at $n_B = 10$) but sacrifices 4.8 points on clean data, i.e. it is *not* no-regret. The gated residual occupies the best clean-vs-robustness trade-off. Notably, the learned gate converges to a nearly constant value (≈ 0.78 ; Fig. 4): rather than “opening” adaptively with blockage, it learns a fixed down-weighting of the position correction that preserves clean accuracy while retaining robustness.

C. Benchmark sensitivity (further analyses)

We characterize the RSRP-only benchmark on clean data. *Width*: increasing the hidden width from 32 to 384 raises top-

TABLE IV
ARCHITECTURE ABLATION: TOP- K ACCURACY ($K=13$, %) VS. BLOCKAGE.

# blocked	0	6	8	10	12
RSRP-only	89.2	70.7	55.5	36.2	22.6
Concatenation	84.4	70.4	60.5	45.5	29.4
Plain residual	88.0	72.1	58.7	40.3	25.0
Gated residual	89.4	72.3	58.7	40.1	25.1




Fig. 4. The learned gate is nearly constant (≈ 0.78) across blockage levels: robustness arises from the always-on residual position path, not from adaptive gating.

K accuracy modestly (89.6% to 91.8% at $K=13$); width 96 is a good accuracy/size balance. *Depth*: 2–4 hidden layers perform equally ($\approx 90\%$), while 6–8 layers degrade slightly (mild overfitting), justifying the four-layer choice. *UE density*: accuracy saturates by $\sim 25\%$ of the training set (3,375 UEs, 90.6%), so the full 15k set is comfortably in the saturated regime—confirming that the clean-data ceiling is not a data-limited artifact.

D. Discussion

Our results delineate the operating regime of multi-modal beam selection. When RSRP is clean and densely informative, single-modality regression is essentially optimal and additional modalities or architectures provide no gain—a useful negative result that tempers the multi-modal intuition for this setting. The position modality becomes valuable only as the RSRP modality degrades, exactly the condition created by mmWave blockage. The gated residual design exploits this asymmetry: it defaults to the RSRP path and adds a bounded position correction, yielding a no-regret method that is strictly preferable to deploying either a position-blind predictor or an always-on (concatenation) fusion that taxes clean-data accuracy. The trade-off exposed in Table IV also offers a design knob: deployments dominated by blockage may prefer concatenation for its larger heavy-blockage gains, accepting a clean-data penalty.

VII. CONCLUSION

We reproduced the SSB-RSRP regression benchmark for 5G NR mmWave beam selection (90% top- K accuracy at $K = 13$) and studied position-aided multi-modal fusion. We found that on clean RSRP the single-modality predictor is at the achievable ceiling, but that a gated residual fusion of sparse RSRP and UE position yields a no-regret robustness improvement: clean-data parity with a consistent 2–4-point top- K accuracy gain under beam blockage. An ablation attributes the no-regret property to the learned gate. Future work includes joint optimization of the measured-beam subset with the fusion network, evaluation on real multi-modal datasets such as DeepSense 6G, and extending the gate to additional impairments (quantization, mobility).

REFERENCES

- [1] Q. Xue, C. Ji, S. Ma, J. Guo, Y. Xu, Q. Chen, and W. Zhang, “A survey of beam management for mmWave and THz communications towards 6G,” *IEEE Commun. Surveys Tuts.*, vol. 26, no. 3, pp. 1520–1559, 2024.
- [2] Y. Heng and J. G. Andrews, “Machine learning-assisted beam alignment for mmWave systems,” in *Proc. IEEE Global Commun. Conf. (GLOBECOM)*, 2019.
- [3] A. Klautau, P. Batista, N. González-Prelcic, Y. Wang, and R. W. Heath, “5G MIMO data for machine learning: Application to beam-selection using deep learning,” in *Proc. Inf. Theory Appl. Workshop (ITA)*, 2018.
- [4] M. S. Sim, Y.-G. Lim, S. H. Park, L. Dai, and C.-B. Chae, “Deep learning-based mmWave beam selection for 5G NR/6G with sub-6 GHz channel information: Algorithms and prototype validation,” *IEEE Access*, vol. 8, pp. 51 634–51 646, 2020.
- [5] M. Polese, F. Restuccia, and T. Melodia, “DeepBeam: Deep waveform learning for coordination-free beam management in mmWave networks,” in *Proc. ACM Int. Symp. Mobile Ad Hoc Netw. Comput. (MobiHoc)*, 2021.
- [6] A. Klautau, N. González-Prelcic, and R. W. Heath, “LiDAR data for deep learning-based mmWave beam-selection,” *IEEE Wireless Commun. Lett.*, vol. 8, no. 3, pp. 909–912, 2019.
- [7] S. Rezaie, E. de Carvalho, and C. N. Manchón, “A deep learning approach to location- and orientation-aided 3D beam selection for mmWave communications,” *IEEE Trans. Wireless Commun.*, vol. 21, no. 12, pp. 11 110–11 124, 2022.
- [8] A. Alkhateeb, G. Charan, T. Osman, A. Hredzak, J. Morais, U. Demirhan, and N. Srinivas, “DeepSense 6G: A large-scale real-world multi-modal sensing and communication dataset,” *IEEE Commun. Mag.*, vol. 61, no. 9, pp. 122–128, 2023.
- [9] U. Demirhan and A. Alkhateeb, “Radar aided 6G beam prediction: Deep learning algorithms and real-world demonstration,” in *Proc. IEEE Wireless Commun. Netw. Conf. (WCNC)*, 2022.
- [10] W. Deng, B. Shi, M. Li, and O. Simeone, “SCAN-BEST: Sub-6GHz-aided near-field beam selection with formal reliability guarantees,” *IEEE Trans. Cogn. Commun. Netw.*, 2025, early access.
- [11] M. Alrabeiah and A. Alkhateeb, “Deep learning for mmWave beam and blockage prediction using sub-6 GHz channels,” *IEEE Trans. Commun.*, vol. 68, no. 9, pp. 5504–5518, 2020.
- [12] S. Jiang, G. Charan, and A. Alkhateeb, “LiDAR aided future beam prediction in real-world millimeter wave V2I communications,” *IEEE Wireless Commun. Lett.*, vol. 12, no. 2, pp. 212–216, 2023.
- [13] A. Alkhateeb and I. Beltagy, “Machine learning for reliable mmWave systems: Blockage prediction and proactive handoff,” in *Proc. IEEE Global Conf. Signal Inf. Process. (GlobalSIP)*, 2018.
- [14] S. Wu, M. Alrabeiah, A. Hredzak, C. Chakrabarti, and A. Alkhateeb, “Deep learning for moving blockage prediction using real mmWave measurements,” in *Proc. IEEE Int. Conf. Commun. (ICC)*, 2022.
- [15] A. Oliveira, A. Khatibi, D. Suzuki, I. Correa, J. Rezende, and A. Klautau, “DL-based beam management for mmWave vehicular networks exploring temporal correlation,” *arXiv preprint arXiv:2511.02260*, 2025.
- [16] 3GPP, “Study on artificial intelligence (AI)/machine learning (ML) for NR air interface,” 3rd Generation Partnership Project (3GPP), Tech. Rep. TR 38.843, 2024.
- [17] —, “Study on channel model for frequencies from 0.5 to 100 GHz,” 3rd Generation Partnership Project (3GPP), Tech. Rep. TR 38.901, 2020.